

# 1 section

## 1.1 subsection

### 1.1.1 subsubsection

- some intermediate title

sample text blabla

math :  $f : X, x = 43, \forall x. \lambda z$

comment

**TODO:** todo

code

Table:

c <sup>1</sup> <sub>4</sub>	c <sup>2</sup> <sub>5</sub>	c <sup>3</sup> <sub>6</sub>
-----------------------------	-----------------------------	-----------------------------

List:

- + pro

- con

Tightcenter:

Text in the center

D.Definition XY here comes the definition

T.Theorem XY here comes a theorem

L.Lemma XY here comes the lemma

Ex.Example XY here comes example

Alg.Algo XY here comes algo

Note: here comes a Note

Important: Here comes something important

Proof: Here comes a proof

Intuition: Here comes some intuition

Raised rule: 

## 2 Maths

'\t test {test}<sup>2</sup> okok

test test {asfa | whenever}<sup>2</sup> okok

test test {test}<sup>2</sup> okok

test test {asfa | whenever}<sup>2</sup> okok

### 3 Introduction

#### 3.1 — Week 1-2: Deep Learning Basics

- MLP
- Fully Connected Networks
- Data, tasks, loss functions
- Backprop
- Activation functions

#### 3.2 — Week 3-4: CNNs, RNNs, co.

- CNNs
- RNNs
- LSTM/GRU/BPTT
- Fully convolutional Networks

#### 3.3 — Week 5-9: Generative Modeling

- Latent variable models
- Implicit Models
- Autoregressive models
- Normalizing Flows/ Invertible Networks

#### 3.4 — Week 10-12: DL 4 CV

- Problems and tasks in human centric computer vision
- DL architectures for CV
- Human body and hand models
- Implicit representations

#### 3.5 — Week 13: Deep RL

**Perceptron:**  $\text{while } \exists x. w^T x > 0 \neq y \text{ do } w = w + \eta(y - \hat{y})x$   
**Note:** If the data is separable, the algorithm converges in finite time

$\hookrightarrow \text{ILP: } \forall l \quad x^{(l)} = \sigma((w^{(l)})^T x^{(l-1)} + b^{(l)}), f(x; w, b) = x^{(L)}$

**Sigmoid:**  $\sigma(x) = \frac{1}{1+e^{-x}} = \frac{e^x}{e^x+1}, \nabla : \sigma(x) \cdot (1 - \sigma(x))$

**Softmax:**  $\text{softmax}(x_i) = \exp(x_i) / \sum_j \exp(x_j)$

**Requirements:** 1) Output must be positive

2) output must be between  $[0, 1]$

3) sum of all outputs must be 1, e.g.  $\sum_i^M \text{softmax}(x_i) = 1$

**TODO:** Linear activation function (Week 2, Page 28)

**MLE:** Maximise  $\log L(\theta) = \log \prod p(x_i|\theta) = \sum \log p(x_i|\theta)$

- 1) Write down probability distribution
- 2) Decompose into per sample probability
- 3) Minimize negative log likelihood

**CE loss:**  $-\frac{1}{N} \sum y_i \log(\sigma(w^T x_i)) + (1 - y_i) \log(1 - \sigma(w^T x_i))$

**Note:** Cross-entropy loss is a maximum likelihood estimator  
We assume the  $y$ 's to be bernoulli distributed, from there we maximise the weights over the probability  $P(D|w) = \prod^N \sigma(w^T x_i)^{y_i} (1 - \sigma(w^T x_i))^{1-y_i}$ . Taking the negative log-likelihood results in the given loss.

**Universal Approximation:**  $\exists g(x) = \sum v_i \sigma(w_i^T x + b_i) \approx f(x)$   
and  $|g(x) - f(x)| < \epsilon$

**Note:**  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$  must be non-constant, bounded and continuous

**Note:**  $f \in [0, 1]^m$ , e.g. must be in the m-dimensional hypercube

## 4 Deep Learning

### 5 Convolutional Neural Network

**Goal:** Good classification performance.

**Trade-off:** Specificity vs. invariance

To achieve this, we need to trade off specificity and invariance (created by affine transformations of an object and different lightings).

The Impact (?) is the generalization ability of the model.

**Receptive Field:** Areas triggering firing of sensory neurons

**Note:** Areas can be on the Retina, Skin, Tongue, etc.

**Note:** Usually divided into excitatory and inhibitory regions.

From Hubel and Wiesels experiments we learned that stimulus covering the whole cat retina, most neurons didn't fire, since the excitatory and inhibitory stimulus canceled each other out. The light must fall on specific regions to excite, forming specific excitement patterns.

**Direction and alignment** of the light affected the firing, and was used to figure out the alignment of the receptive field.

**Hierarchy:** Simple cell respond to simple 0-1 input (noisy).

Complex cells, connected to multiple simple cells, form more complex patterns, more resistant to invariance.

**Note:** Size of receptive fields tends to get larger, the more complex the cell becomes.

Sensory receptors connected to cell in the brain form the receptive field of that cell. This receptive field can be very noisy in simple cells.

The connection of such simple cells to complex cells can form complex triggering patterns, f.e. a line in a certain direction. This hierarchy can be extended to hypercomplex shapes.

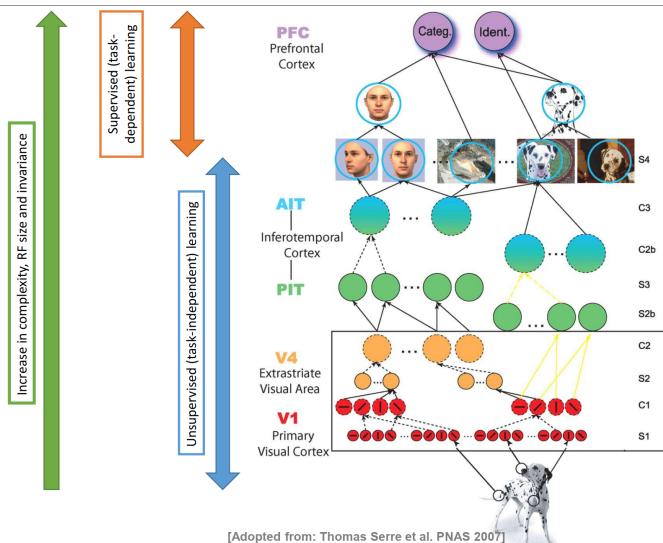


Figure 1: Complexity, unsupervised and supervised learning hierarchy

**HMAX:**  $S: y = \exp(-\frac{1}{2\sigma^2} \sum_{j=1}^{n_{S_k}} (w_j - x_j)^2)$ ,  $C: y = \max_{n_{C_k}} y_j$

$x$  is the input to the cell,  $w$  is the weight of the cell.  $\sigma$  tells

us about the size of the receptive fields.

**Conv:** Modify pixel by some function of surrounding pixels

**Note:** Any linear, shift-equivariant transform can be written as convolution.

**Linear:**  $T(\alpha u + \beta v) = \alpha T(u) + \beta T(v)$

**Invariant:**  $T(f(u)) = f(T(u))$

**Note:** We want this in classification, b.c. a cat in the middle of the picture should still be classified as cat if it is shifted or rotated

**Equivariant:**  $T(f(u)) = f(T(u))$

**Note:** In Edge detection very important, if we shift the edge in the image, we also want the response to shift the same way

**Lin. Filtering:**  $I'(i, j) = \sum_{(m,n) \in N(i,j)} K(m, n) I(i+m, j+n)$

**Note:**  $I$  is the image,  $K$  is the Kernel,  $N(i,j)$  is the neighbourhood of a pixel.

**Shift-invariance:** The Kernel is usually parameterized as  $K(i,j,m,n)$ , e.g. the weights of the Kernel depend on the location in the image. Removing  $i,j$ -dependence makes the kernel shift invariant.

1

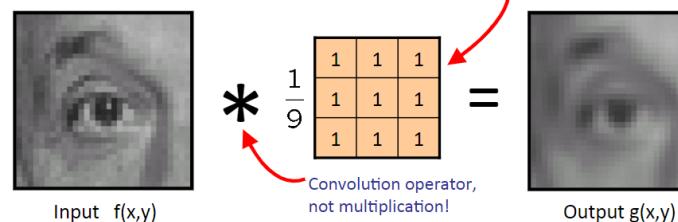


Figure 2: Linear filtering is applied with the convolution operator to compute a new image from the neighbourhood of each pixel of the old image

**Correlation:**  $I'(i, j) = \sum_{m=-k}^k \sum_{n=-k}^k K(m, n) I(i+m, j+n)$

**Conv:**  $I'(i, j) = \sum_{m=-k}^k \sum_{n=-k}^k K(m, n) I(i-m, j-n)$

Alternative representation would be

$I'(i, j) = \sum_{m=-k}^k \sum_{n=-k}^k K(-m, -n) I(i+m, j+n)$

**Note:** Convolution can be done via matrix multiplication. It acts as a point-spread function, what we want is the weights of this kernel

**TODO:** Ask how good we must know this, e.g. perform it ourselves?

$\nabla$ : Conv. with kernel  $[-1, 1]$ ,  $\frac{\delta f}{\delta x} \approx \frac{f(x_{n+1}, y) - f(x_n, y)}{\Delta x}$

**Correlation = Convolution:** iff  $K(i, j) = K(-i, -j)$

**Convolutional Layer:**  $w^T z^{(l-1)} + b$  (+ activation function)

$I'$  dim:  $\frac{I_{height} + 2 \cdot \text{padding} - \text{dilation} \cdot (K_{height} - 1) - 1}{\text{stride}} + 1$

**Note:** If we have 6 filters, the produced "new image" would be of size 28x28x6

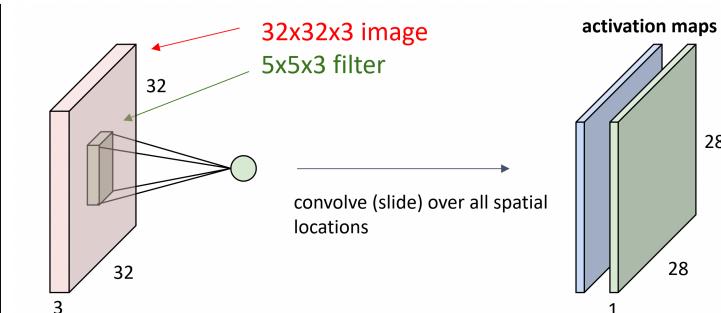


Figure 3: Visualisation applying filters to an image

**TODO:** Do convolution via matrix on your own!

**Weight sharing:** Same feature detector on whole image and ↓ weights

**Note:** This makes the feature detector robust against affine transformations, as it detects the feature in the whole image once trained

**Stride:** used to reduce size, replaces pooling layers

**Dilation:** add "holes" to filter, fast increase of rec. field

It is easy to integrate the global context like this

**CNN-fwd:**  $z_{i,j}^{(l)} = w^{(l)} * z^{(l-1)} + b = (\sum_{m,n} w_{m,n}^{(l)} z_{i-m, j-n}^{(l-1)}) + b$

$$\begin{aligned} I'(i', j') &= w * I(i, j) \\ &= \sum_{m,n} w_{m,n} I(i-m, j-n) \end{aligned}$$

Figure 4: Visualisation forward pass

**CNN-bwd (z):**  $\delta_{i,j}^{(l-1)} = \frac{\delta C}{\delta z_{i,j}^{(l-1)}} = \sum_{i',j'} \frac{\delta C}{\delta z_{i',j'}^{(l)}} \frac{\delta z_{i',j'}^{(l)}}{\delta z_{i,j}^{(l-1)}}$

$= \sum_{i',j'} \delta_{i',j'}^{(l)} w_{i'-i, j'-j}^{(l)}$

If we take the derivative of  $\delta z_{i',j'}^{(l)}$ , only certain weight terms stick, the rest is going to be zero.

**Note:** Backward path is just a convolution with the flipped kernel

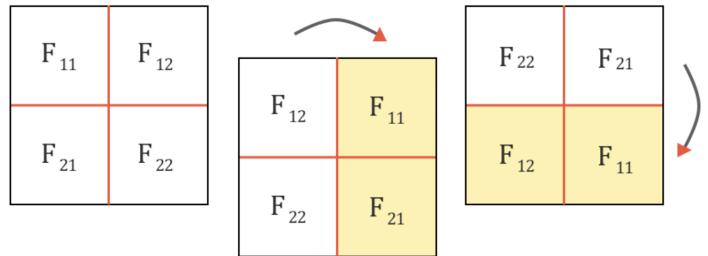


Figure 5: Flipped convolution operator

**TODO:** Do calculation yourself to see why it is flipped, or better, understand the relation to the "standart" convolution

CNN-bwd (w):  $\delta_w^{(l)} * ROT_{180}(z^{(l-1)})$

**QUESTION:** Is this going to be added or substracted to the original image?

**Note:** Each layer holds more complex "patterns", this resembles the neuron complexity hierarchy we have seen in nature.

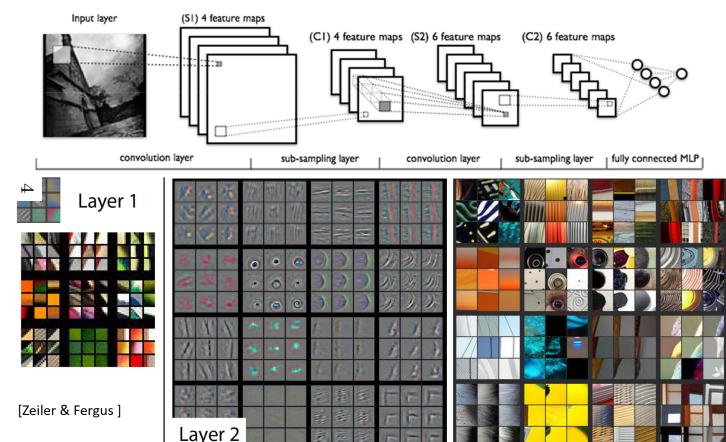


Figure 6: Visualisation of output of CNN layers

**Max Pooling:**  $m \times m$  Filter with  $m$  stride, take max. Reduces size of each activation layer, downsampling. This helps in extending the activation region of downstream pixels, as we are putting more info in a smaller region. It also helps in reducing the noise in pictures.

There are other pooling strategies that can be applied, but with the MAX-pooling, we gain robustness to local changes. For example, if a digit rotates a bit, the region probably still has the same max pixel and thus is robust against rotations.

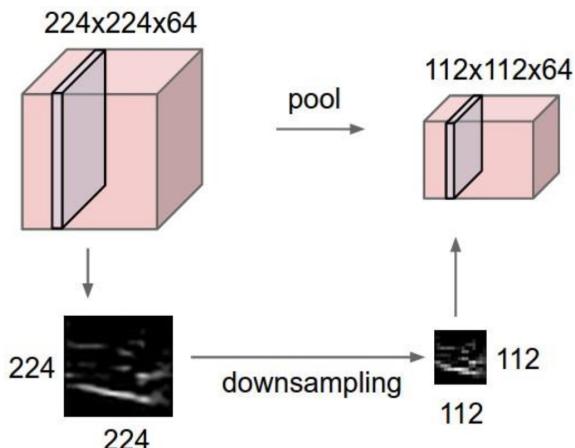


Figure 7: Visualisation of pooling layer

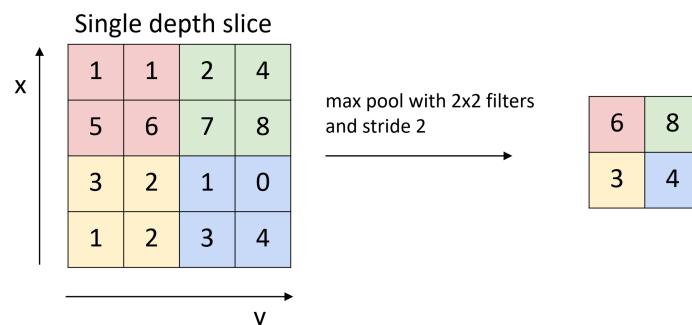


Figure 8: Visualisation of max-pooling layer

fwd:  $z^{(l)} = \max\{z_i^{(l-1)}\}$  bwd:  $\frac{\delta z^{(l)}}{\delta z_i^{(l-1)}} = 1 \text{ if } i^* = \max\{z_i^{(l-1)}\}$

Only one contributing pixel gets a gradient, all others have gradient zero.

## 5.1 — Evolution of architectures

### Revolution of Depth

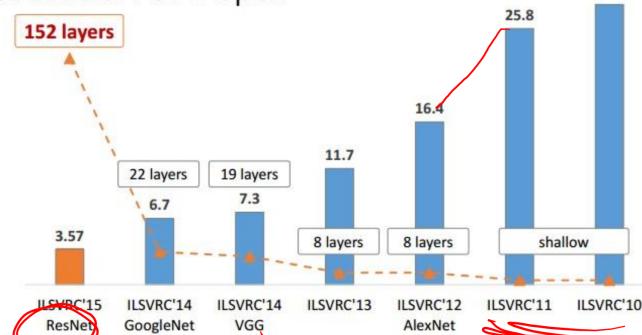


Figure 9: Evolution of NN architecture

## 5.2 — VGG vs. AlexNet

less kernels/filters ( $\downarrow$  params), more layers ( $\uparrow$  perceptive field)  
Larger receptive field means that the net can respond to patterns that are larger spread apart.  
Each large filter, f.e. 11x11, can be represented using multiple 3x3 filters. The number of parameters for one 11x11 filter is 121, which is equal to using 5 3x3 filters with 45 parameters.

## 5.3 — GoogleNet

More layers, removed fully connected layer on the top.

**Inception:** use 1x1 conv. layers to reduce layer depth

Because of small kernels and deep networks, the number of channels (depth) is going to be huge. Inception modules down-samples the activation maps, reducing the number of channels by convolution, and thus reducing the number of parameters.

**Note:** In addition, GoogLeNet uses auxiliary classification heads throughout the CNN, to make sure the gradients don't dry out.

## 5.4 — ResNet

**Residual-Con:** Skips weight layers with residual connections.

**Note:** A deep network should at least perform as good as a shallow one (set all additional layers to identity).

This approach lead to the believe that failing to perform is an optimisation issue, not a design issue.

ResNet one (drastic) downsampling and then residual layers. The "Ultra-deep" version is not used often, in general ResNet18 is used.

### • Residual net

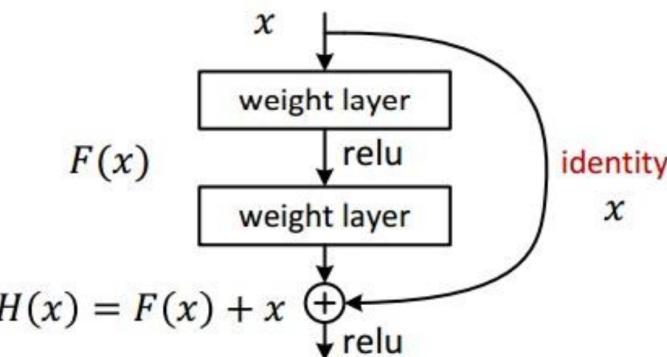


Figure 10: Evolution of NN architecture

## 6 Fully Convolutional Neural Network

The main difference to classical CNN is that we don't want to only work with fixed sized images as input. This limitation comes from the fully connected layers in the end.

**Goal:** Pixel-to-pixel classification By removing the size constraints of the input pictures allows us to classify pixels by using their local neighbourhood.

**Drawback:** We don't get the information of the whole context. The basic approach is using a standard CNN to learn features of the input image (encoding), which is followed by a decoder who projects the learned features to the higher resolution pixel space.

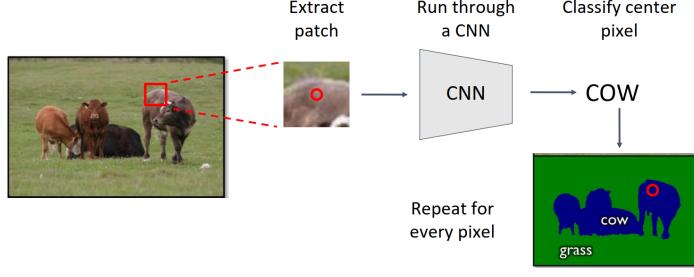


Figure 11: Pixel-wise classification pipeline

**Low-res:** due to all the pooling and down-sampling

This results in low resolution and fuzzy object boundaries. Deeper architectures can mitigate this by upsampling again. The learned features on each downsampled level are copied to the upsampling (decoding) stream, to give the fcn a chance to include features of different levels.

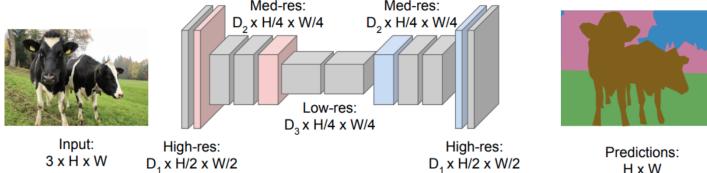
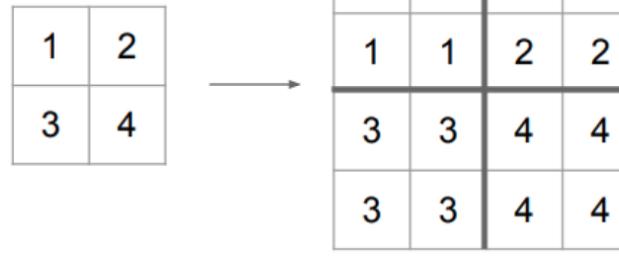


Figure 12: Downsample and then upsample to increase resolution again

**Upsampling (NN):** copy value for the whole output

Upsampling can be seen as interpolation, increasing the resolution of the signal.

## Nearest Neighbor



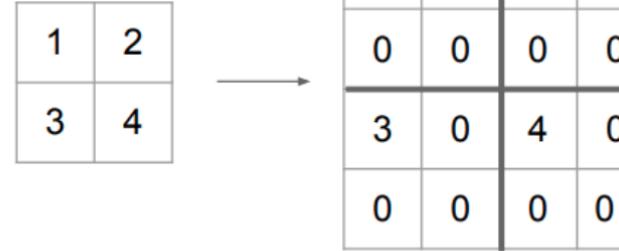
Input: 2 x 2

Output: 4 x 4

Figure 13: Copy input pixel to all output pixels

**Bed-of-nails:** zero all outputs but one copy of input

## “Bed of Nails”



Input: 2 x 2

Output: 4 x 4

Figure 14: Copy input pixel to one output pixels, rest zero

**Max-Unpooling:** remember max pooling, BOF to that location

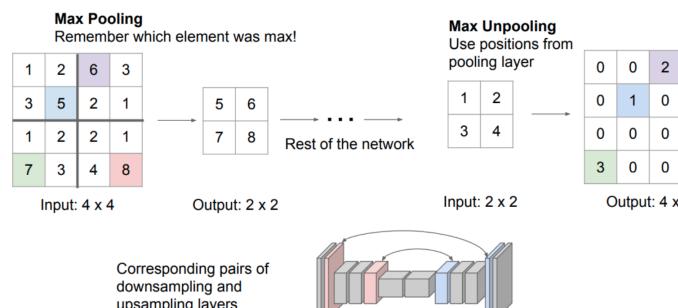


Figure 15: Copy input pixel to location of the max pooling pixel

**Transpose Conv:** input gives (learnable) weight for filter

3 x 3 transpose convolution, stride 2 pad 1

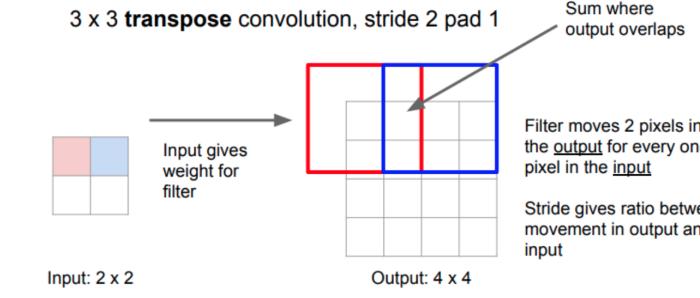


Figure 16: Input pixel gives the weight for a filter

**UNet:** TransCov + Skip connections

Main idea: combine global and local feature maps by copying corresponding tensors from earlier stages.

Also, the downsampling information is pretty shallow, upsampling performance is improved when including “pre-downsampling” information.

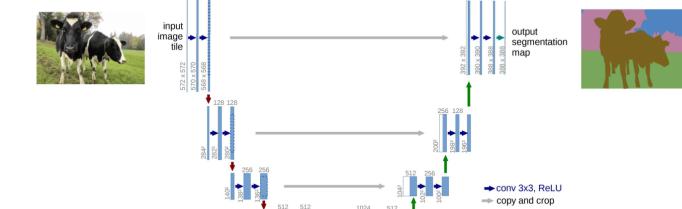


Figure 17: U-Net architecture

## 7 Recurrent Neural Network

The main advantage of RNN's over vanilla NN's or CNN's is the flexibility when it comes to input size.

NN has a fixed input/output size and computes a fixed amount of computational steps (defined by number of layers).

RNN can work over sequences of data, learning the transition given the (implicitly stored) past data.

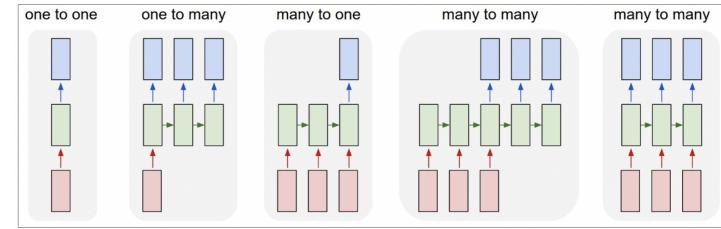


Figure 18: RNN 1. NN, 2. Image captioning, 3. Sentiment analysis, 4. Translation, 5. Video classification

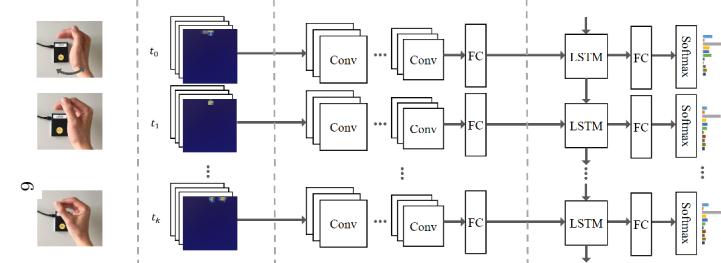


Figure 19: RNN Video classification architecture

**Vanilla RNN:**  $\hat{y} = W_{hy}h^t$ ,  $h^t = \tanh(W_{hh}h^{t-1} + W_{xh}x^t)$

**Goal:** Learn transition function knowing what's worth keeping. Tanh is important, otherwise we can't forget information.

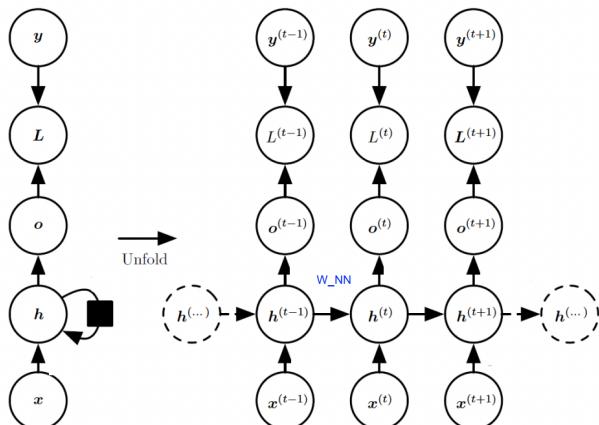


Figure 20: Visualisation of complex RNN

**Backprop:**  $\frac{\delta L}{\delta W} = \sum_{t=1}^S \frac{\delta L^t}{\delta W} = \sum_{t=1}^S \sum_{k=1}^t \frac{\delta L^t}{\delta y^t} \frac{\delta y^t}{\delta h^t} \frac{\delta h^t}{\delta h^k} \frac{\delta h^k}{\delta W}$

**TODO:** Understand this better

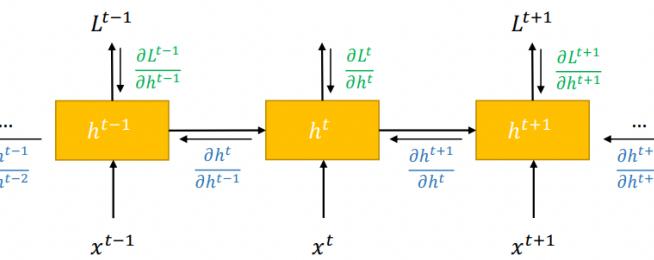


Figure 21: RNN Backprop unrolled

**Gradient:**  $h^t = W^T h^{t-1} = (W^T)^t h^1 = (Q^T \Lambda^t Q) h^1$

The Eigenvalues of the weigh matrix explode or vanish when the sequence length gets to big. Regularizing the Eigenvalues is proven to reduce the learning capabilities of the learning-system drastically, and is thus not an option

**Exploding gradients:** exploding gradients can be managed by clipping after a certain size

**Vanishing gradients:** Add memory cell

$$\frac{\partial L^t}{\partial W} = \sum_{k=1}^t \frac{\partial L^t}{\partial y^t} \frac{\partial y^t}{\partial h^t} \frac{\partial h^t}{\partial h^k} \frac{\partial h^k}{\partial W}$$

we are interested in this

$$\frac{\partial h^t}{\partial h^k} = \prod_{i=k+1}^t \frac{\partial h^i}{\partial h^{i-1}} = \prod_{i=k+1}^t W_{hh}^T \text{diag}[f'(h^{i-1})]$$

$$\forall i, \left\| \frac{\partial h^i}{\partial h^{i-1}} \right\| \leq \|W_{hh}^T\| \|\text{diag}[f'(h^{i-1})]\| < \frac{1}{\gamma} \gamma < 1$$

$$\left\| \frac{\partial h^t}{\partial h^k} \right\| < (\eta)^{t-k}$$

Figure 22: Proof that vanishing Gradients are a problem

**LSTM:**  $c_t^l = f \odot c_{t-1}^l + i \odot g$ ,  $h_t = o \odot \tanh(c_t^l)$ ,

$f, i, o, g = [\sigma, \sigma, \sigma, \tanh] \odot W^l [h_{t-1}^l h^{l-1}]^T$

To mitigate the vanishing gradients, we introduce a long term memory unit. There are three phases to this: forgetting, new information and output generation.

- 1) Forgetting happens through manipulating  $c$  with  $f \in [0, 1]$
- 2) New information is compiled by trying to figure out what to add/ remove to the current memory, as  $g \in [-1, 1]$
- 3) Output generation is done by mixing the short-term input with the long-term memory

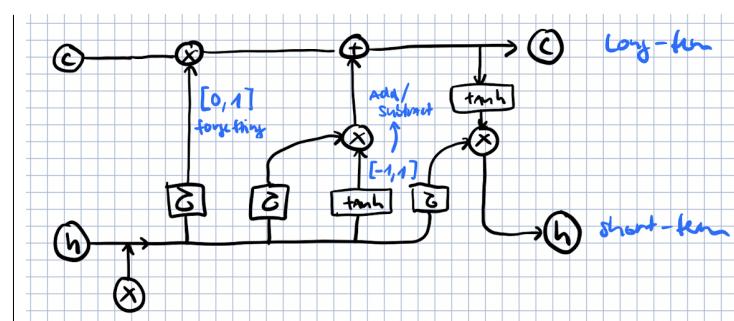


Figure 23: LSTM Visualization

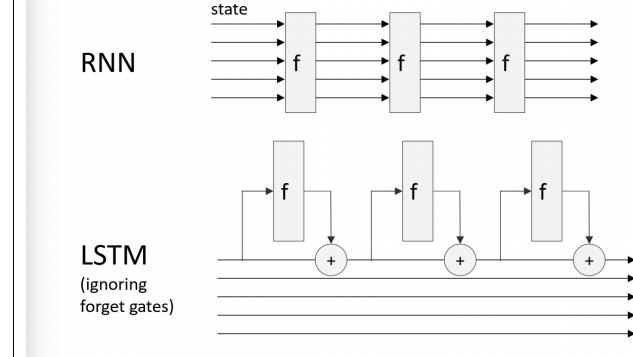


Figure 24: LSTM vs. RNN

## 8 Variational Auto Encoders

Supervised vs. Unsupervised learning:

**Supervised learning:** Tries to learn a function that maps  $X \Rightarrow Y$ , used in classification, regression, object detection and segmentation.

**Unsupervised Learning:** The Goal here is to learn the hidden structure of the data, such as in clustering, feature learning, dimensionality reduction or density estimation.

**Generative Modeling (our Goal):** Given training data, learn a distribution and generate new samples drawn from the learned distribution. Learn  $p_{model}$  close to  $p_{data}$ , generate from  $p_{model}$

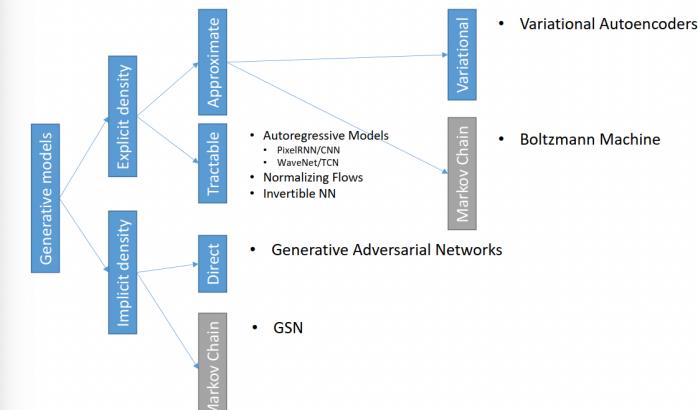


Figure 25: Taxonomy of variational models

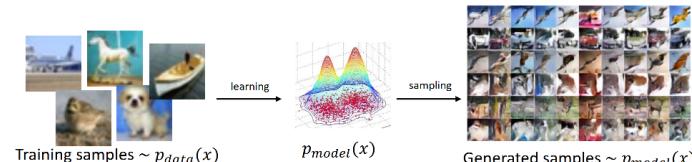


Figure 26: VAE Pipeline

$$\text{AE: optimize } \theta_f, \theta_g = \arg \min \sum_n^N \|x_n - g_\theta(f_\theta(x_n))\|^2$$

**The encoder** projects the original input to a latent space  $Z$ . It has to figure out what are the important features worth keeping (f.e. when  $\dim(Z) < \dim(X)$ ).

**The decoder** learns to construct an output from a sample of  $Z$ .

In the optimal case, the encoder-decoder pipeline approximates the identity function of the data.

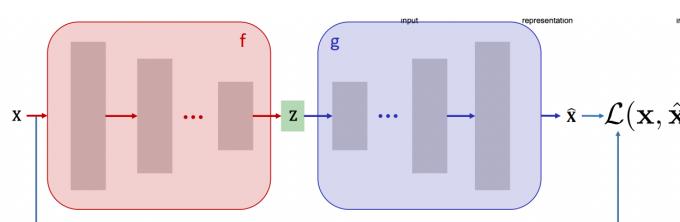


Figure 27: Auto-Encoders Model with loss

**PCA:**  $z = f(x) = Wx + b$ ,  $\hat{x} = g(z) = W^*z + c$   
linear embedding along the principle components.

**NN:**  $z = f(x) = \sigma(Wx + b)$ ,  $\hat{x} = g(\hat{a}(x)) = \sigma(W^*z + c)$

**Latent Space:** meaningful DOF, optimally continuous and interpolatable, under- or overcomplete

Not all DOF are important, when we look at the image space (f.e.  $256 \times 256 \times 3$ ), we could as well sample something very random which does not represent an image.

**Undercomplete Z:** Compressed representation, learns important features of input. Bad for out-of-distribution samples

**Overcomplete Z:** Hidden units copies input components, but might not extract meaningful structure.

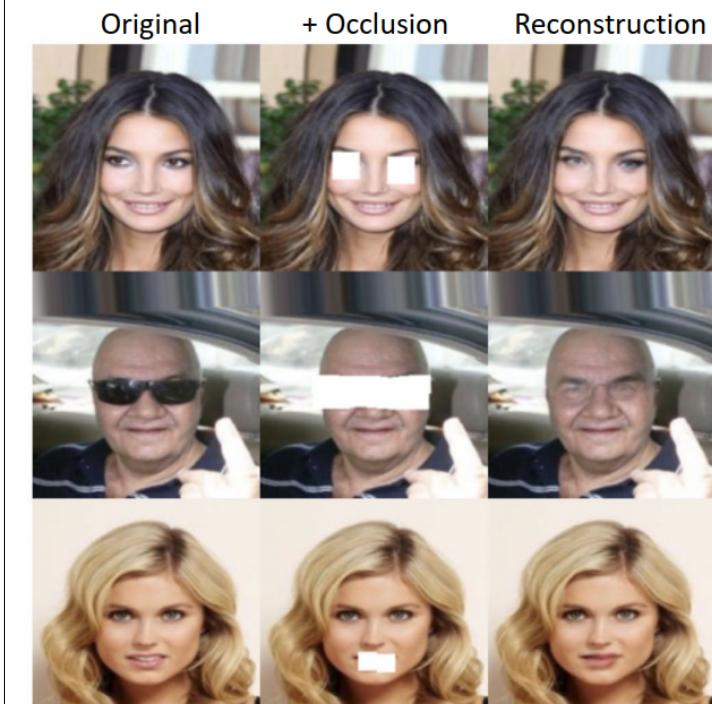


Figure 30: VAE de-noising example

**Limit:** latent space not interpolatable, unseen samples?  
Freshly generated samples from latent space  $Z$  do not make much sense, as in these regions, the decoder doesn't know what makes sense. Still, for reconstruction, the latent space is very good.

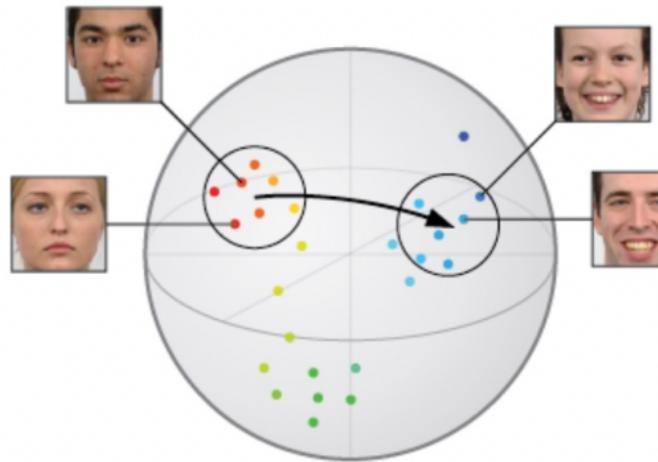


Figure 28: Latent space

**Denoising:** add gaussian noise to input, reconstruct to original

**Note:** The Latent space is overcomplet in this instance

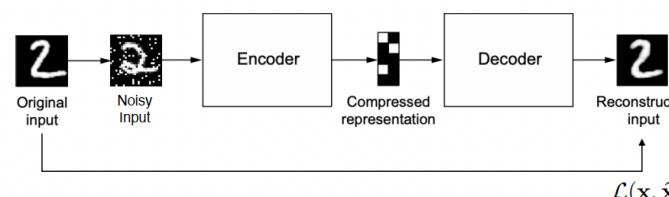


Figure 29: VAE pipeline for denoising

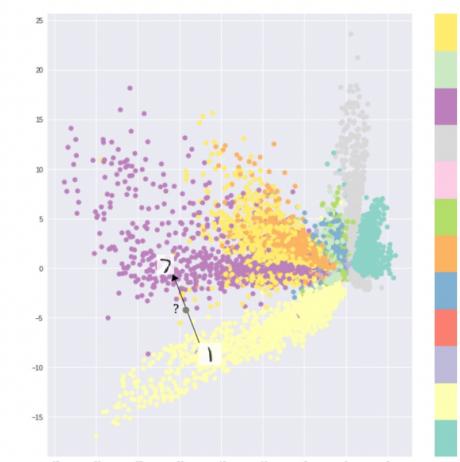


Figure 31: AE MNIST embedding, points in the space

VAE: The latent space Z is approximated by  $f(x) \sim \mathcal{N}(\hat{\mu}, \hat{\sigma}I)$

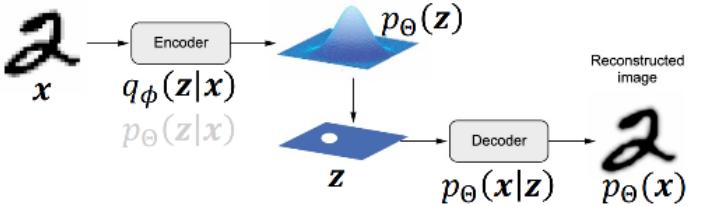


Figure 32: VAE Learned distributions

By design, the latent space is compact and interpolatable. This allows easy sampling and interpolation when generating new samples.

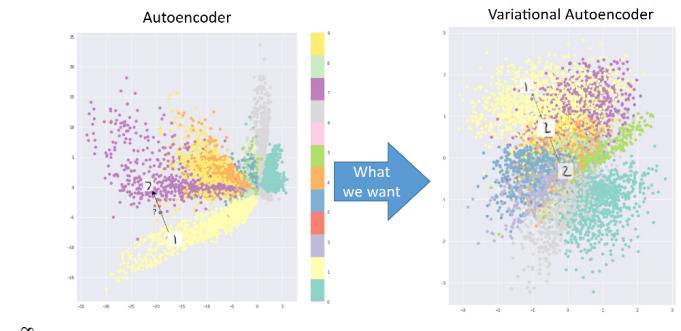


Figure 33: VAE Latent space

**Encode:**  $q_\phi(z|x)$  for  $\mu_{(z|x)}, \Sigma_{(z|x)}$ ,  $z|x \sim \mathcal{N}(\mu_{(z|x)}, \Sigma_{(z|x)})$

The Encoder outputs a vector each for  $\mu, \sigma$ . This multivariate gaussian can then be sampled to get the latent vector Z depending on x.

**Problem:** Using only the reconstruction error (MSE), the variance is going to be reduced to zero and the individual classes form clusters around their means. This leads to bad interpolation performance. We want to force the encoder to find features that are as close as possible to each other, while still being distinct.

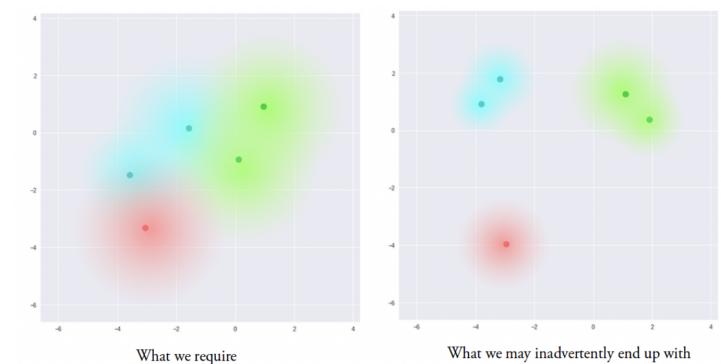


Figure 34: VAE Reconstruction loss only

**Decode:**  $p_\theta(x|z)$  for  $\mu_{(\hat{x}|z)}, \Sigma_{(\hat{x}|z)}$ ,  $\hat{x}|z \sim \mathcal{N}(\mu_{(\hat{x}|z)}, \Sigma_{(\hat{x}|z)})$

$$\begin{aligned} \text{KL: } & -D_{KL}(q_\phi(z|x)||p_\theta(z)) = \int_x q_\phi(z|x) \log\left(\frac{p_\theta(z)}{q_\phi(z|x)}\right) \\ & = \frac{1}{2} \sum_i^J (1 + \log(\sigma_j^2) - \mu_j - \sigma_j^2), \text{ if } q \sim \mathcal{N}(\mu, \sigma I), p \sim \mathcal{N}(0, I) \\ \text{from exercise: } & \int p(z) \log q(z) dz = \\ & -\frac{J}{2} \log 2\pi - \frac{1}{2} \sum_i^J \log \sigma_{p,j} - \frac{1}{2} \sum_j \frac{\sigma_{p,j}^2 + (\mu_{p,j} - \mu_{q,j})^2}{\sigma_{q,j}^2} \end{aligned}$$

KL measures how similar two distributions are. We punish the encoder if the final distribution diverges away from the standard normal. **Note:** Assymmetric and non-negative

**ELBO:**  $p_\theta(x) \geq E_{q_\phi(z|x)}[\log p_\theta(x|z)] - D_{KL}(q_\phi(z|x)||p_\theta(z))$

The first part of the loss is the reconstruction loss. It measures the loss from the decoder prediction.

The second part of the loss is the latent code loss. It enforces that the latent space is a gaussian, e.g. enforcing properties such as compactness and good interpolation

**Note:** We want to maximize the ELBO

**Note:** If we have a sequence, the ELBO is the sum of all variational lower bounds

**Note:** The different classes are still entangled, interpolation works good, but features are not distinguishing unique properties.

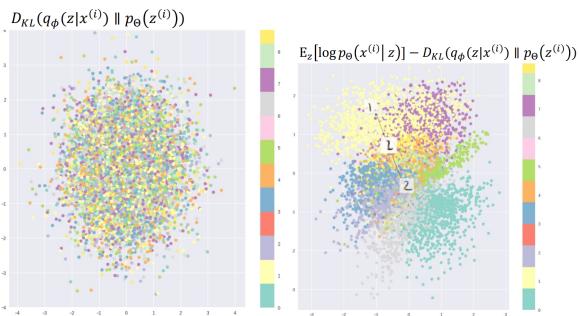


Figure 35: VAE KL-Loss only vs. both losses (entangled)

**Goal:** Learn features that correspond to distinct factors of variation e.g., digits and style (or thickness, orientation) **Statistical independence of the features can be used**  
In the picture below, one can see that the AE might achieve disentangled features, but it is bad for interpolation. The KL divergence loss creates some entanglement, though.

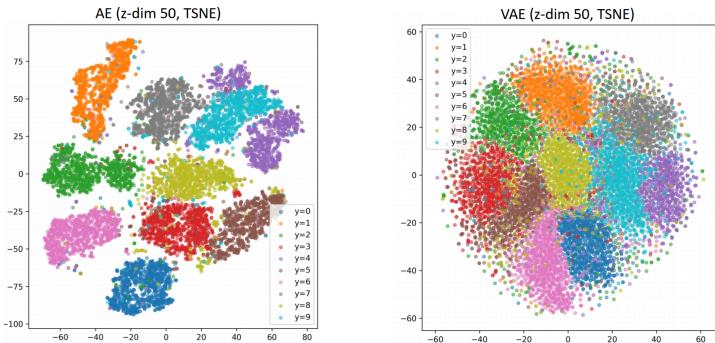


Figure 36: VAE vs. AE

**$\beta$ -VAE:**  $\mathcal{L} = -E_{q_\phi(z|x)}[\log p_\theta(x|z)] + \beta D_{KL}(q_\phi(z|x)||p_\theta(z))$

**Goal** is to disentangle the different features. The KL-loss tries to create a diagonal multivariate Gaussian, which perfectly separates the independent features. The reconstruction loss fucks this up, as it fights for the loss as well.

Putting more weight for the KL loss helps solving it in an unsupervised fashion. Alternatively, could condition on features.

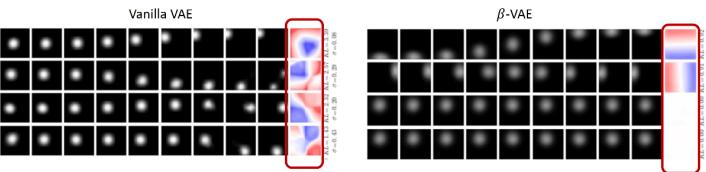


Figure 37: VAE Entangled vs. Disentangled representation

## 9 Auto-regressive Models

**Regressive Property:**  $x_t = b_0 + b_1 x_{t-1} + b_2 x_{t-2}$

**Sequence Model:**  $p(x) = \prod^N p(x_i|x_1, \dots, x_{i-1}) = \prod p(x_i|x_{<i})$

We are trying to estimate a distribution here.

**Note:** Sequence models trivially fulfill the regressive property

**Believe Net:**  $\hat{x}_i = p(x_i = 1|x_{<i}) = Ber(\sigma(\sum_1^{i-1} w_j^i x_j + w_0^i))$

In this fully visible sigmoid belief network, the value of the next element is modelled as a Bernoulli variable.

**Note:** The weight Matrix  $W \in \mathbb{R}^{n \times n}$  is triangular.

Element  $i$  at row  $j$  contains  $w_j^i$ , describing weight of  $x_i$  when calculating  $\hat{x}_j$ .

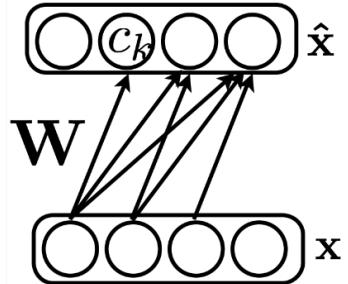


Figure 38: AR Fully Visible Sigmoid Belief Network

⇒ **NADE:**  $h_i = \sigma(b + \sum^{i-1} w_j x_j)$  (in  $O(H)$ , update in  $O(1)$ )  
 $\hat{x}_i = p(x_i = 1|x_{<i}) = \sigma(c_i + V_i h_i)$  ( $N$  times ⇒ total  $O(NH)$ )  
The difference to FVSB-Net is that the weights are shared.  
Therefore, the number of weights is in  $\mathcal{O}(n)$ .

**Note:**  $b + W_{<i+1}x_{<i+1} = b + W_{<i}x_{<i} + W_{i+1}x_{i+1}$  in  $\mathcal{O}(1)$

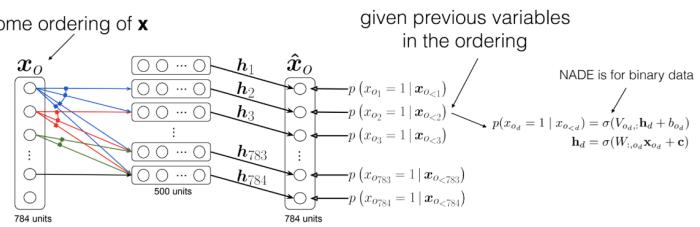


Figure 39: AR NADE Schema

**PixelRNN:**  $p(x) = \prod^{N^2} p(x_i|x_{<i}) = p(x_{i,R}|x_{<i})p(x_{i,G}|x_{<i}, x_{i,R})$   
RNN's are autoreg. from nature:  $h_t$  summarises inputs  $< t$   
Learning and generation is slow, because of the explicit pixel dependencies from the pixel ordering.  
If we use a LSTM for example,  $h_{i,j} = f(h_{i-1,j}, h_{i,j-1}, p_{i,j})$ , e.g. no parallelization in training, neither in prediction, can be made. It is very slow, but LSTM and RNN can capture long-term dependencies very well.

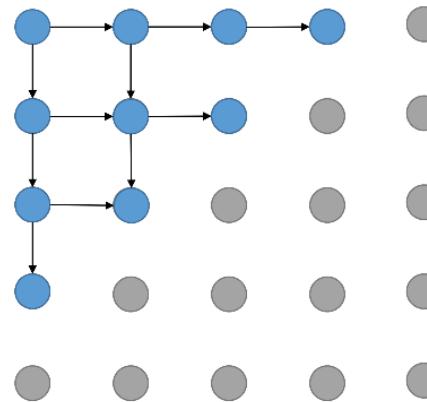


Figure 40: AR PixelRNN dependencies

**PixelCNN:** Use conv + mask, parallelize training, blind spot

**Note:** Generation is still sequential

**Summary:**

Pros are the explicit likelihood  $p(x)$  and good evaluation metric through LL of training data.

Cons are slow generation

**WaveNet:** dilated convolution for large scale temp. dependencies

The issue with sound is that the frequency is at  $\sim 16'000$  samples/second.

To capture speech, convolution must have a very large receptive field.

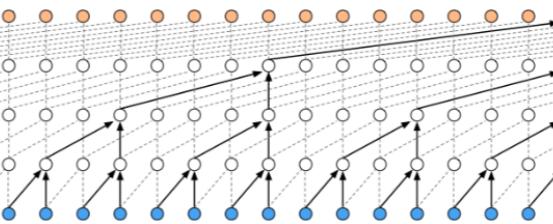


Figure 41: AR WaveNet long term dependencies

**Dilated Conv:** exponential receptive field size increase

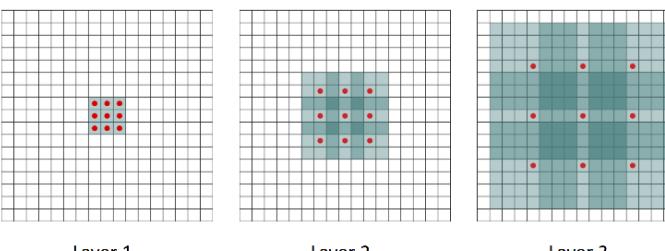


Figure 42: Dilated Conv: Receptive field increase 1

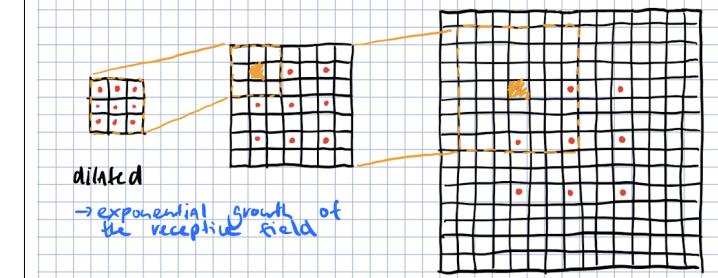


Figure 43: Dilated Conv: Receptive field increase 2

**Self-Attention:**  $K = XW_K, V = XW_V, Q = x_t W_Q (\in \mathbb{R}^D), X \in \mathbb{R}^{T \times D}, W \in \mathbb{R}^{D \times D}, \alpha = softmax(QK^T / \sqrt{D}), x_{t+1} = \alpha V X = softmax(\frac{(W_Q X)(W_K X)^T}{\sqrt{D}} + M)(W_V X)$

**RNN:**  $h_t = g(x_{t-1}, h_{t-1}), x_t = f(x_{t-1}, h_t)$ . Learns to encode relevant information for future steps and maintains long-term temporal relations.

Complexity per layer:  $\mathcal{O}(TD^2)$ , with  $\mathcal{O}(T)$  sequential operations

**CNN:**  $x_t = f(x_{t-1}, x_{t-2}, \dots, x_1)$ . Learns to aggregate past information for the next step.

Complexity per layer:  $\mathcal{O}(kTD^2)$ , with  $\mathcal{O}(1)$  sequential operations

**Self Attention:**  $x_t = f(x_{t-1}, x_{t-2}, \dots, x_1)$ . Learns to identify relevant information for the next step.

Very expressive, but computationally complex. Attention is a quadratic operation. Complexity per layer:  $\mathcal{O}(T^2D)$ , with  $\mathcal{O}(1)$  sequential operations. Can be restricted in how far we look back, reduces complexity of layer

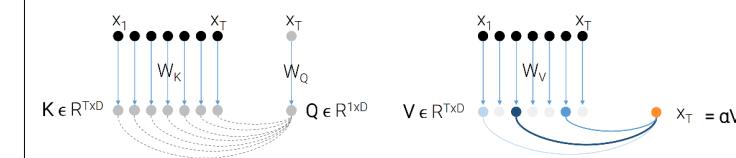


Figure 44: AR Self-attention

**TODO:** Understand the runtimes better, be able to derive this

## 10 Normalizing Flows

**Goal:** What we want is a model that combines the tractable likelihoods from AR models with the latent space of VAE.

**Variable change 1D:**  $p_x(x) = p_z(h(x)) |h(x)|$

**2D:**  $\int \int f(x, y) dx dy = \int \int f(g(u, v), h(u, v)) J(u, v) du dv$

The probabilistic mass must be preserved!

**Normalizing Flow:**  $f : \mathbb{R} \rightarrow \mathbb{R}$ , cont. and invertible

$$p_X(x; \theta) = p_Z(f_\theta^{-1}(x)) \left| \det\left(\frac{\partial f_\theta^{-1}(x)}{\partial x}\right) \right| = p_Z(f^{-1}(x)) \left| \det\left(\frac{\partial f(x)}{\partial z}\right) \right|^{-1}$$

**Example Linear transformation:**  $z = Ax$

$$\Rightarrow p_x(x) = p_z(A^{-1}x) |\det(A^{-1})|$$

**Note:** Not every NN works to model f:

- must be differentiable
- must be invertible
- must preserve dimensionality
- jacobian must be computed efficiently

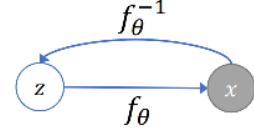


Figure 45: Normalizing Flow mapping

**Coupling:**  $(y_A; y_B)^T = (h(x^A, \beta(x^B)); x^B)^T \mid h: \text{element}, \beta: \text{NN}$

$$(x^A; x^B)^T = (h^{-1}(y^A; \beta(y^B)); y^B)^T, J = ((h'; h'f'), (0; 1))$$

**Note:** The Jacobian is triangular. This allows computation of  $\det$  in  $\mathcal{O}(d)$ , instead of  $\mathcal{O}(d^3)$

**Note:**  $\beta$  does not have to be invertible

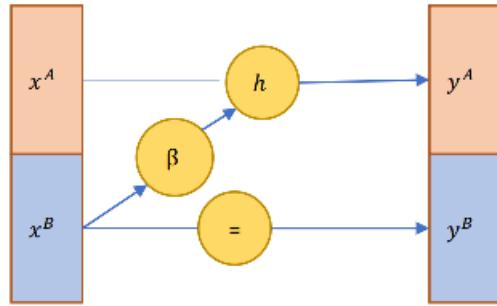


Figure 46: Normalizing Flow coupling layer

**Composition:**  $p_X(x; \theta) = p_Z(f_\theta^{-1}(x)) \prod_k \left| \det\left(\frac{\partial f_\theta^{-1}(x)}{\partial x}\right) \right|$

For  $x = f(z) = f_k \circ f_{k-1} \circ \dots \circ f_2 \circ f_1(z)$

**Train:**  $\log p_x(D) = \sum_x^D (\log p_z(f^{-1}(x)) + \sum_k \log \left| \det\left(\frac{\partial f^{-1}(x)}{\partial x}\right) \right|)$

**Inference:** draw  $z \sim p_z(\cdot)$ ,  $\hat{x} = f(z)$

The whole point of this is that we want to draw  $z \sim \mathcal{N}(0, I)$  and then learn the transformation to create really complex x

**Model:**  $x_i \rightarrow \text{squeeze} \rightarrow n.\text{flow} \rightarrow \text{split} \rightarrow z_i$  (repeat  $L - 1$ )

$\rightarrow \text{squeeze} \rightarrow n.\text{flow} \rightarrow z_L$

The flow step is repeated K times and consists of

- actnorm (like batch normalization)
- 1x1 conv
- affine coupling layer